

University POLITEHNICA of Bucharest

Faculty of Automatic Control and Computers,
Computer Science and Engineering Department



Diploma

Minimizing Dockerfiles

Scientific Adviser:
Prof. Răzvan Deaconescu

Author:
Andrei Petrea

Bucharest, 2025

I want first and foremost to thank my advisor, Prof. Răzvan Deaconescu, for his guidance and support throughout this project. His expertise and insights have been invaluable in shaping my understanding of the subject matter and in helping me navigate the challenges I faced during this endeavor.

I also want to thank my friends and family without whom these 4 years would have been so much harder.

Contents

Acknowledgements	i
1 Introduction	1
2 Background	2
2.1 A brief history of containers	2
2.2 The anatomy of a container	2
2.3 Related work	3
3 Motivation and Objectives	4
4 Use Cases	5
5 Building Blocks	6
6 Architecture Overview	7
6.1 Static Analysis	7
6.2 Dynamic Analysis	7
6.3 Brute Force	7
7 Implementation Details	9
7.1 Overview	9
7.2 Project Structure	9
7.3 Preprocessing	10
7.4 Static Analysis	10
7.5 Dynamic Analysis	12
7.6 Binary Search	13
8 Performance Evaluation	15
9 Status and Planned Work	16
9.1 Status	16
9.2 Planned Work	16

List of Figures

2.1	Container layers	2
4.1	High-level comparison of the software components of traditional VMs (a), containers (b), containers in VMs (c) with unikernels (d).	5
5.1	High-level architecture of the application	6
6.1	Brute force approach to finding the minimal Dockerfile	8
7.1	Language execution speed comparison [1]	9
8.1	Minimal dockerfile comparision	15

List of Listings

2.1	Sample Dockerfile	2
2.2	Sample Multi-stage Dockerfile	2
4.1	Sample Kraftfile	5
6.1	ldd command	7
6.2	strace command	7
7.1	Project Structure	9
7.2	Go DockerConfig structure	10
7.3	Dockerfile.minimal.template	10
7.4	ldd parser function	10
7.5	File util functions	11
7.6	File util functions	12
7.7	File util functions	12
7.8	Binary search	13

Chapter 1

Introduction

Chapter 2

Background

2.1 A brief history of containers

2.2 The anatomy of a container

```
1 FROM python:3.10.0-slim
2 COPY requirements.txt .
3 RUN pip install -r requirements.txt
4 COPY . .
5 CMD ["python", "adapter.py"]
```

Listing 2.1: Sample Dockerfile

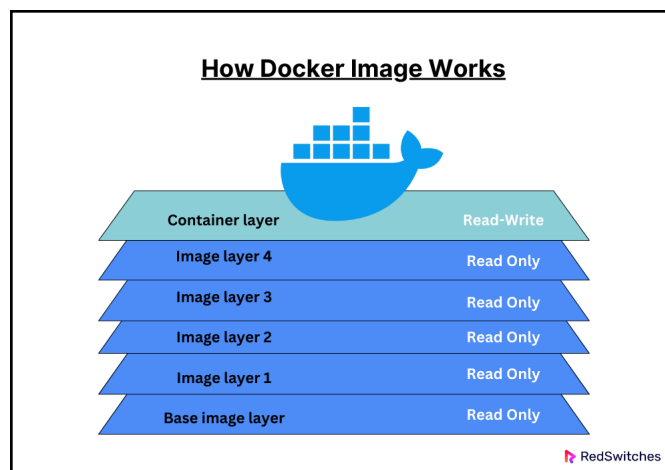


Figure 2.1: Container layers

```
1 FROM alpine:latest AS builder
2 RUN apk --no-cache add build-base
3
4 FROM builder AS build1
5 COPY source1.cpp source.cpp
6 RUN g++ -o /binary source.cpp
7
```

```
8  FROM builder AS build2
9  COPY source2.cpp source.cpp
10 RUN g++ -o /binary source.cpp
```

Listing 2.2: Sample Multi-stage Dockerfile

2.3 Related work

Chapter 3

Motivation and Objectives

Chapter 4

Use Cases

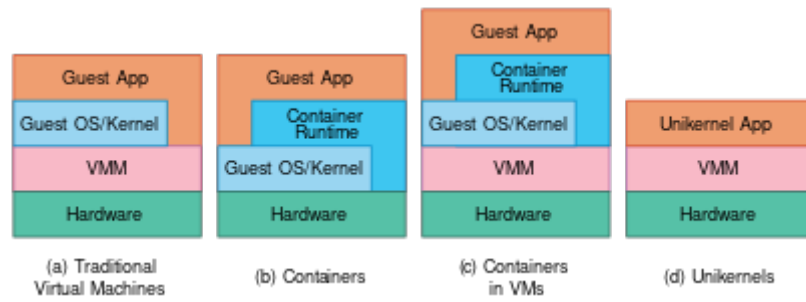


Figure 4.1: High-level comparison of the software components of traditional VMs (a), containers (b), containers in VMs (c) with unikernels (d).

```
1 spec: v0.6
2 runtime: base:latest
3 rootfs: ./Dockerfile
4 cmd: ["/helloworld"]
```

Listing 4.1: Sample Kraftfile

Chapter 5

Building Blocks

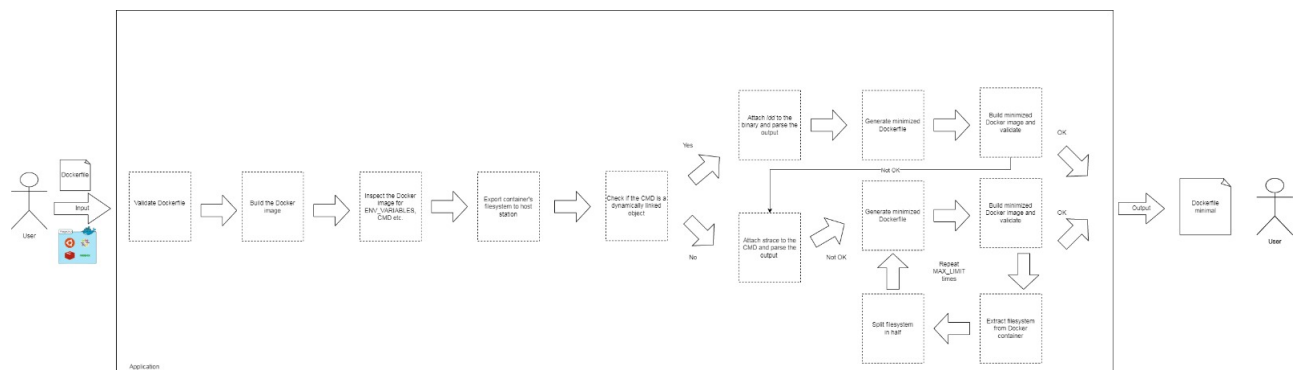


Figure 5.1: High-level architecture of the application

Chapter 6

Architecture Overview

6.1 Static Analysis

```
1 ~ > ldd /usr/bin/man
2     linux-vdso.so.1 (0x00007ffe831f5000)
3     libmandb-2.9.1.so => /usr/lib/man-db/libmandb-2.9.1.so (0
4         x00007f068b572000)
5     libman-2.9.1.so => /usr/lib/man-db/libman-2.9.1.so (0
6         x00007f068b52f000)
7     libz.so.1 => /lib/x86_64-linux-gnu/libz.so.1 (0x00007f068b502000)
8     libpipeline.so.1 => /lib/x86_64-linux-gnu/libpipeline.so.1 (0
9         x00007f068b4f1000)
10    libc.so.6 => /lib/x86_64-linux-gnu/libc.so.6 (0x00007f068b2ff000)
11    libgdbm.so.6 => /lib/x86_64-linux-gnu/libgdbm.so.6 (0
12        x00007f068b2ef000)
13    libseccomp.so.2 => /lib/x86_64-linux-gnu/libseccomp.so.2 (0
14        x00007f068b2cb000)
15    /lib64/ld-linux-x86-64.so.2 (0x00007f068b59b000)
```

Listing 6.1: ldd command

6.2 Dynamic Analysis

```
1 ~ > strace -fe execve,openat echo "Hello, World\!"
2     execve("/usr/bin/echo", ["echo", "Hello, World!"], 0x7fff43b04ce8 /* 36
3         vars */) = 0
4     openat(AT_FDCWD, "/etc/ld.so.cache", O_RDONLY|O_CLOEXEC) = 3
5     openat(AT_FDCWD, "/lib/x86_64-linux-gnu/libc.so.6", O_RDONLY|O_CLOEXEC)
6         = 3
7     Hello, World!
8     +++ exited with 0 +++
```

Listing 6.2: strace command

6.3 Brute Force

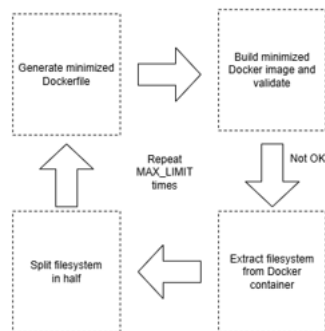


Figure 6.1: Brute force approach to finding the minimal Dockerfile

Chapter 7

Implementation Details

7.1 Overview

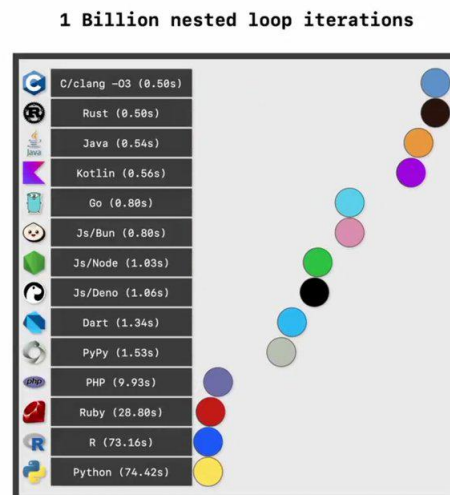


Figure 7.1: Language execution speed comparison [1]

7.2 Project Structure

```
1      .
2      |-- cmd
3      |-- main.go
4      |-- dockerminimizer
5      |-- dockerminimizer.go
6      |-- go.mod
7      |-- go.sum
8      |-- install.sh
9      |-- ldd
10     |-- ldd.go
11     |-- logger
12     |-- logger.go
13     |-- preprocess
```

```

14     |-- preprocess.go
15     |-- README.md
16     |-- strace
17     |-- strace.go
18     |-- types
19     |-- types.go
20     |-- utils
21     |-- utils.go

```

Listing 7.1: Project Structure

7.3 Preprocessing

```

1     type DockerConfig struct {
2         User          string          `json:"User"`
3         ExposedPorts   map[string]map[string]any `json:"ExposedPorts"`
4         Env            []string        `json:"Env"`
5         Cmd            []string        `json:"Cmd"`
6         WorkingDir     string          `json:"WorkingDir"`
7         Entrypoint     []string        `json:"Entrypoint"`
8     }

```

Listing 7.2: Go DockerConfig structure

```

1 FROM node:20.9.0 as builder
2
3 WORKDIR /app
4
5 COPY app.js /app/app.js
6 COPY package.json /app/package.json
7 COPY package-lock.json /app/package-lock.json
8
9 RUN npm install
10 EXPOSE 3000
11 CMD ["node", "app.js"]
12
13
14 FROM scratch
15
16 ENV PATH=/usr/local/sbin:/usr/local/bin:/usr/sbin:/usr/bin:/sbin:/bin
17 ENV NODE_VERSION=20.9.0
18 ENV YARN_VERSION=1.22.19
19 WORKDIR /app
20 EXPOSE 3000/tcp
21 ENTRYPOINT ["docker-entrypoint.sh"]
22 CMD ["node", "app.js"]

```

Listing 7.3: Dockerfile.minimal.template

7.4 Static Analysis

```

1     func ParseOutput(output []byte, rootfsPath string) (map[string][]string,
2         map[string]string) {
3         files := make(map[string][]string)
4         symLinks := make(map[string]string)
5         scanner := bufio.NewScanner(bytes.NewReader(output))

```

```

5      for scanner.Scan() {
6          line := scanner.Text()
7          if strings.Contains(line, ">") {
8              parts := strings.Split(line, ">")
9              lib := strings.Split(strings.TrimSpace(parts[1]), " ")[0]
10
11              if strings.Contains(lib, "not found") {
12                  continue
13              }
14              utils.AddFilesToDockerfile(lib, files, symLinks, rootfsPath)
15
16              } else if strings.Contains(line, "not found") {
17                  continue
18              } else {
19                  lib := strings.Split(strings.TrimSpace(line), " ")[0]
20                  utils.AddFilesToDockerfile(lib, files, symLinks, rootfsPath)
21              }
22          }
23      return files, symLinks
24  }
```

Listing 7.4: ldd parser function

```

1  func RealPath(path string) string {
2      realPath, _ := filepath.Abs(path)
3      return filepath.Clean(realPath)
4  }
5
6  func CheckIfFileExists(file string, envPath string) bool {
7      info, err := os.Stat(envPath + "/" + file)
8      return !os.IsNotExist(err) && !info.IsDir()
9  }
10
11 func CheckIfSymbolicLink(file string, envPath string) bool {
12     info, err := os.Lstat(envPath + "/" + file)
13     if err != nil {
14         return false
15     }
16     return info.Mode()&os.ModeSymlink != 0
17 }
18
19 func ReadSymbolicLink(file string, envPath string) string {
20     link, _ := os.Readlink(envPath + "/" + file)
21     resolved := link
22     if !filepath.IsAbs(link) {
23         resolved = filepath.Join(filepath.Dir(envPath+"/"+file),
24                                 link)
25     }
26     return strings.TrimPrefix(resolved, envPath)
27 }
28
29 func AddFilesToDockerfile(file string, files map[string][]string, symLinks
30 map[string]string, rootfsPath string) {
31     file = RealPath(file)
32     if CheckIfFileExists(file, rootfsPath) {
33         if CheckIfSymbolicLink(file, rootfsPath) {
34             symLinks[file] = ReadSymbolicLink(file, rootfsPath)
35         } else {

```



```

34             files[filepath.Dir(file)] = AppendIfMissing(files[
35                 filepath.Dir(file)], file)
36         }
37     }

```

Listing 7.5: File util functions

7.5 Dynamic Analysis

For this stage, we attach a *strace* probe to the docker running command, in order to determine the files used and processes that are spawned by the command. In order to do this, we utilize Docker's volumes to mount a statically-linked *strace* binary into the container, as well as a logfile to save the output of the command. The reason for using a statically-linked binary is so that the libraries used by *strace* don't interfere with the libraries used by the command. To determine the runtime dependencies of the command, we will trace the *open* and *exec* family of syscalls [2] as well as passing the *-f* flag to trace the child processes as well.

```

1  func parseOutput(output string, syscalls []string, files map[string][]
2      string, symLinks map[string]string, envPath string) {
3      regexes := make(map[string]*regexp.Regexp)
4      for _, syscall := range syscalls {
5          regexes[syscall] = regexp.MustCompile(syscall + `
6              \([^\)]*?\("[^"]+"\)`
7      )
8      for line := range strings.SplitSeq(output, "\n") {
9          for _, syscall := range syscalls {
10             if regexes[syscall].MatchString(line) {
11                 match := regexes[syscall].FindStringSubmatch(line)
12                 if len(match) > 1 {
13                     utils.AddFilesToDockerfile(match[1], files, symLinks
14                         , envPath+"/rootfs")
15                 }
16                 break
17             }
18         }
19     }
20 }

```

Listing 7.6: File util functions

```

1  func getSheBang(command string, rootfsPath string) string {
2      file, err := os.Open(rootfsPath + "/" + command)
3      if err != nil {
4          log.Error("Failed to open file:", command)
5          return ""
6      }
7      defer file.Close()
8      var firstLine []byte
9      buf := make([]byte, 1)
10     for {
11         n, err := file.Read(buf)
12         if n > 0 {
13             if buf[0] == '\n' {
14                 break
15             }
16             firstLine = append(firstLine, buf[0])

```

```

17         }
18         if err != nil {
19             break
20         }
21     }
22     return string(firstLine)
23 }
24
25 func parseShebang(imageName string, containerName string, syscalls []string,
26     files map[string][]string, symLinks map[string]string, envPath
27     string, metadata types.DockerConfig, timeout int) (map[string][]
28     string, map[string]string) {
29     command := utils.GetContainerCommand(imageName, envPath, metadata)
30     hasSudo := utils.HasSudo()
31     shebang := getSheBang(command, envPath+"/rootfs")
32     regex := regexp.MustCompile(`^#!\s*([^\s]+)`)
33     if !regex.MatchString(shebang) {
34         log.Error("Failed to find shebang in file:", command)
35         return files, symLinks
36     }
37     match := regex.FindStringSubmatch(shebang)
38     if len(match) < 2 {
39         log.Error("Failed to find interpreter in shebang:", command)
40         return files, symLinks
41     }
42     interpreter := match[1]
43     lddCommand := hasSudo + " chroot " + envPath + "/rootfs ldd " +
44         interpreter
45     log.Info("Running command:", lddCommand)
46     lddOutput, err := exec.Command("sh", "-c", lddCommand).
47         CombinedOutput()
48     if err != nil {
49         log.Error("Failed to run ldd command\n" + err.Error())
50     }
51     files, symLinks = ldd.ParseOutput(lddOutput, envPath+"/rootfs")
52
53     output := getStraceOutput(imageName, envPath+"/strace", envPath+"/
54         log.txt", syscalls,
55         containerName, interpreter, envPath, metadata, timeout)
56     parseOutput(output, syscalls, files, symLinks, envPath)
57     return files, symLinks
58 }

```

Listing 7.7: File util functions

7.6 Binary Search

```

1 func binarySearchStep(envPath string, context string, timeout int, step int,
2     usedFiles map[string][]string,
3     unusedFiles map[string][]string) (map[string][]string, map[string][]
4     string, error) {
5     for {
6         if len(unusedFiles) == 0 {
7             return nil, nil, errors.New("no unused files or
8                 symbolic links left to process")
9         }
10    }

```

```

7         usedFiles, unusedFiles = splitFilesystem(usedFiles,
            unusedFiles)
8         filename := fmt.Sprintf("Dockerfile.minimal.binary_search.%d",
            step)
9         usedFiles = utils.ShrinkDictionary(usedFiles)
10        utils.CreateDockerfile(filename, "Dockerfile.minimal.
            template", envPath, usedFiles, nil)
11        err := utils.ValidateDockerfile(filename, envPath, context,
            timeout)
12        if err == nil {
13            return make(map[string][]string), usedFiles, nil
14        }
15    }
16 }
17
18 func BinarySearch(envPath string, maxLimit int, context string, timeout int)
    error {
19     log.Info("Starting binary search...")
20     usedFiles, unusedFiles, err := parseFilesystem(envPath + "/rootfs")
21     if err != nil {
22         log.Error("Error parsing filesystem:", err)
23         return errors.New("error parsing filesystem")
24     }
25     step := 1
26     for i := range maxLimit {
27         log.Info("Binary search iteration:", step)
28         usedFiles, unusedFiles, err = binarySearchStep(envPath,
            context, timeout, step,
29             usedFiles, unusedFiles)
30         if err != nil {
31             break
32         }
33         step = i + 1
34     }
35     if step == maxLimit {
36         log.Info("Reached maximum limit of binary search iterations:
            ", maxLimit)
37         return errors.New("reached maximum limit of binary search
            iterations")
38     }
39     log.Info("Binary search completed sucessfully.")
40     return nil
41 }

```

Listing 7.8: Binary search

Chapter 8

Performance Evaluation

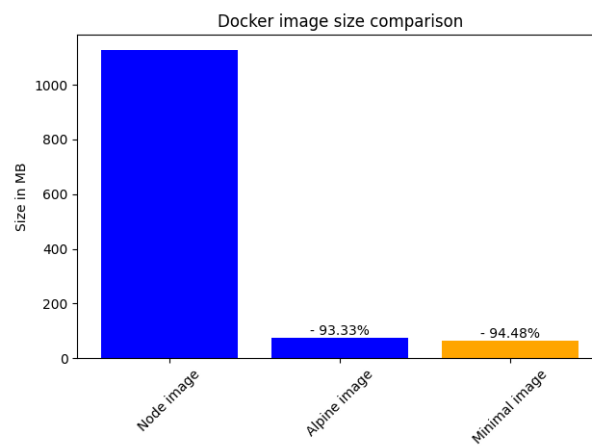


Figure 8.1: Minimal dockerfile comparision

Chapter 9

Status and Planned Work

9.1 Status

9.2 Planned Work

Bibliography

- [1] Ben Dicken. Speed comparison of programming languages. <https://benjdd.com/languages/>.
- [2] Adam Rehn. Identifying application runtime dependencies. <https://unrealcontainers.com/blog/identifying-application-runtime-dependencies>.